

Prediction and Control with Function Approximation: Learning Objectives

Module 00: Welcome to the Course

Understand the prerequisites, goals and roadmap for the course.

Module 01: On-policy prediction with approximation

Lesson 1: Estimating Value Functions as Supervised Learning

Understand how we can use parameterized functions to approximate value functions

Explain the meaning of linear value function approximation

Recognize that the tabular case is a special case of linear value function approximation.

Understand that there are many ways to parameterize an approximate value function

Understand what is meant by generalization and discrimination

Understand how generalization can be beneficial

Explain why we want both generalization and discrimination from our function approximation

Understand how value estimation can be framed as a supervised learning problem

Recognize not all function approximation methods are well suited for reinforcement learning

Lesson 2: The Objective for On-policy Prediction

Understand the mean-squared value error objective for policy evaluation

Explain the role of the state distribution in the objective

Understand the idea behind gradient descent and stochastic gradient descent

Outline the gradient Monte Carlo algorithm for value estimation

Understand how state aggregation can be used to approximate the value function

Apply Gradient Monte-Carlo with state aggregation

Lesson 3: The Objective for TD

Understand the TD-update for function approximation

Highlight the advantages of TD compared to Monte-Carlo

Outline the Semi-gradient TD(0) algorithm for value estimation

Understand that TD converges to a biased value estimate

Understand that TD converges much faster than Gradient Monte Carlo

Lesson 4: Linear TD

Derive the TD-update with linear function approximation

Understand that tabular TD(0) is a special case of linear semi-gradient TD(0)

Highlight the advantages of linear value function approximation over nonlinear

Understand the fixed point of linear TD learning

Describe a theoretical guarantee on the mean squared value error at the TD fixed point

Module 02: Constructing Features for Prediction

Lesson 1: Feature Construction for Linear Methods

Describe the difference between coarse coding and tabular representations

Explain the trade-off when designing representations between discrimination and generalization

Understand how different coarse coding schemes affect the functions that can be represented

Explain how tile coding is a (computationally?) convenient case of coarse coding

Describe how designing the tilings affects the resultant representation

Understand that tile coding is a computationally efficient implementation of coarse coding

Lesson 2: Neural Networks

Define a neural network.

Define activation functions.

Define a feedforward architecture

Understand how neural networks are doing feature construction

Understand how neural networks are a non-linear function of state

Understand how deep networks are a composition of layers.

Understand the tradeoff between learning capacity and challenges presented by deeper networks

Lesson 3: Training Neural Networks

Compute the gradient of a single hidden layer neural network

Understand how to compute the gradient for arbitrarily deep networks

Understand the importance of initialization for neural networks

Describe strategies for initializing neural networks

Describe optimization techniques for training neural networks

Module 03: Control with Approximation

Lesson 1: Episodic Sarsa with Function Approximation

Explain the update for Episodic Sarsa with function approximation

Introduce the feature choices, including passing actions to features or stacking state features

Visualize value function and learning curves

Discuss how this extends to Q-learning easily, since its a subset of Expected Sarsa

Lesson 2: Exploration under Function Approximation

Understanding optimistically initializing your value function as a form of exploration

Lesson 3: Average Reward

Describe the average reward setting

Explain when average reward optimal policies are different from discounted solutions

Understand how differential value functions are different from discounted value functions

Module 04: Policy Gradient

Lesson 1: Learning Parameterized Policies

Understand how to define policies as parameterized functions

Define one class of parameterized policies based on the softmax function

Understand the advantages of using parameterized policies over action-value based methods

Lesson 2: Policy Gradient for Continuing Tasks

Describe the objective for policy gradient algorithms

Describe the results of the policy gradient theorem

Understand the importance of the policy gradient theorem

Lesson 3: Actor-Critic for Continuing Tasks

Derive a sample-based estimate for the gradient of the average reward objective

Describe the actor-critic algorithm for control with function approximation, for continuing tasks

Lesson 4: Policy Parameterizations

Derive the actor-critic update for a softmax policy with linear action preferences

Implement this algorithm

Design concrete function approximators for an average reward actor-critic algorithm

Analyze the performance of an average reward agent

Derive the actor-critic update for a gaussian policy

Apply average reward actor-critic with a gaussian policy to a particular task with continuous actions

